

Introduction to Online Learning Algorithms

Yoav Freund

January 6, 2025

Outline

About this Course

Halving Algorithm

Perceptron

Estimating the mean

Class web site

- ▶ All of the class material is available from the github repository

<https://github.com/yoavfreund/2025-online-learning>



What the class will cover

- ▶ Introduction (with Mean)
- ▶ Exponential weights algorithms
 - ▶ Hedge
 - ▶ Mixability
 - ▶ Bregman Divergences
- ▶ Online learning and Coding
 - ▶ Universal Coding
 - ▶ Continuous Experts
 - ▶ The Context Algorithm
- ▶ Multiple arm Bandit
- ▶ Tracking
 - ▶ Tracking
 - ▶ Tracking within a small set of experts
- ▶ Online learning and game theory
 - ▶ Repeated Matrix Games
 - ▶ Internal regret.
 - ▶ Drifting games
 - ▶ NormalHedge
- ▶ Online Convex Optimizatio
 - ▶ Follow the regularized leader
 - ▶ Dual Descent
 - ▶ AdaGrad

HW / Evaluation

- ▶ 5 HW assignments for $5 \cdot 15 = 75$ opints
- ▶ A final for 25 points.


Example trace for Halving Algorithm

	$t = 1$	$t = 2$	$t = 3$	$t = 4$	$t = 5$
expert1	1	1	1	1	-
expert2	1	0	-	-	-
expert3	0	-	-	-	-
expert4	1	0	-	-	-
expert5	1	0	-	-	-
expert6	0	-	-	-	-
expert7	1	1	1	1	-
expert8	1	1	1	0	0
alg.	1	0	1	1	0
outcome	1	1	1	0	0

Mistake bound for Halving algorithm

- ▶ Each time algorithm makes a mistake, the pool of perfect experts is halved (at least).
- ▶ We assume that at least one expert is perfect.
- ▶ Number of mistakes is at most $\log_2 N$.
- ▶ No stochastic assumptions whatsoever.
- ▶ Proof is based on combining a lower and upper bounds on the number of perfect experts.

The Perceptron Problem




- ▶ $\|\vec{V}\| = 1$
- ▶ Example = (\vec{X}, y) ,
 $y \in \{-1, +1\}$.
- ▶ $\forall \vec{X}, \|\vec{X}\| \leq R$.
- ▶ $\forall (\vec{X}, y),$
 $y(\vec{X} \cdot \vec{V}) \geq g$

The Perceptron learning algorithm

- ▶ An online algorithm. Examples presented one by one.
- ▶ start with $\vec{W}_0 = \vec{0}$.
- ▶ If mistake: $(\vec{W}_i \cdot \vec{X}_i)y_i \leq 0$
 - ▶ Update $\vec{W}_{i+1} = \vec{W}_i + y_i X_i$.

Example trace for the perceptron algorithm




Bound on number of mistakes

- ▶ The number of mistakes that the perceptron algorithm can make is at most $\left(\frac{R}{g}\right)^2$.
- ▶ Proof by combining upper and lower bounds on $\|\vec{W}\|$.

Pythagorean Lemma

If $(\vec{W}_i \cdot \vec{X}_i)y < 0$ then

$$\|\vec{W}_{i+1}\|^2 = \|\vec{W}_i + y_i \vec{X}_i\|^2 \leq \|\vec{W}_i\|^2 + \|\vec{X}_i\|^2$$



Upper bound on $\|\vec{W}_i\|$

Proof by induction

- ▶ Claim: $\|\vec{W}_i\|^2 \leq iR^2$
- ▶ Base: $i = 0$, $\|\vec{W}_0\|^2 = 0$
- ▶ Induction step (assume for i and prove for $i + 1$):
$$\begin{aligned}\|\vec{W}_{i+1}\|^2 &\leq \|\vec{W}_i\|^2 + \|\vec{X}_i\|^2 \\ &\leq \|\vec{W}_i\|^2 + R^2 \leq (i + 1)R^2\end{aligned}$$

Lower bound on $\|\vec{W}_i\|$

$\|\vec{W}_i\| \geq \vec{W}_i \cdot \vec{V}$ because $\|\vec{V}\| = 1$.

Let i denote the number of mistakes made so far.

We prove a lower bound on $\vec{W}_i \cdot \vec{V}$ by induction over i

- ▶ Claim: $\vec{W}_i \cdot \vec{V} \geq ig$
- ▶ Base: $i = 0$, $\vec{W}_0 \cdot \vec{V} = 0$
- ▶ Induction step (assume for i and prove for $i + 1$):
$$\begin{aligned}\vec{W}_{i+1} \cdot \vec{V} &= (\vec{W}_i + \vec{X}_i y_i) \cdot \vec{V} = \vec{W}_i \cdot \vec{V} + y_i \vec{X}_i \cdot \vec{V} \\ &\geq ig + g = (i + 1)g\end{aligned}$$

Combining the upper and lower bounds

$$(ig)^2 \leq \|\vec{W}_i\|^2 \leq iR^2$$

Thus:

$$i \leq \left(\frac{R}{g}\right)^2$$

The mean estimation game

- ▶ An adversary chooses a real number $y_t \in [0, 1]$ and keeps it secret.
- ▶ You make a guess of the secret number x_t
- ▶ The adversary reveals the secret and you pay $(x_t - y_t)^2$
- ▶ You want to minimize $\frac{1}{T} \sum_{t=1}^T (x_t - y_t)^2$
- ▶ Impossible without additional constraints.

Adversary is a fixed distribution

- ▶ Suppose that the adversary draws y_1, y_2, \dots, y_T IID from a fixed distribution over $[0, 1]$ with mean μ and std σ .
- ▶ Optimal prediction $x_t = \mu$
- ▶ $E_Y [(\mu - Y)^2] = \sigma^2$
- ▶ Online prediction: predict x_{t+1} from $Y^t = \langle Y_1, Y_2, \dots, Y_t \rangle$.
- ▶ **Expected regret:** compare performance of algorithm to $\text{Regret} = E_{Y^T} [(x_t - Y_t)^2] - \sigma^2$

Individual sequence bounds

- ▶ Make no assumption about how the sequence is generated.
- ▶ The best constant value for x in hind-sight:

$$x_T^* \doteq \operatorname{argmin}_{x \in [0,1]} \sum_{t=1}^T (x - y_t)^2, \quad x_t^* = \frac{1}{T} \sum_{t=1}^T y_t$$

- ▶ Regret: the loss over and above the loss of x_T^* . **for the worst-case sequence**

$$\operatorname{Regret}_T = \sum_{t=1}^T (x_t - y_t)^2 - \sum_{t=1}^T (x_T^* - y_t)^2$$

- ▶ **Goal:** sublinear regret $\lim_{T \rightarrow \infty} \frac{\operatorname{Regret}_T}{T} = 0$

Follow the Leader

- ▶ Idea: set x_{t+1} to be the best constant prediction on y_1, \dots, y_t
- ▶ $x_{t+1} = \operatorname{argmin}_{x \in [0,1]} \sum_{i=1}^t (x - y_i)^2$
- ▶ We will prove that the regret of this algorithm is upper bound by $4 + 4 \ln T$

regret bound

Theorem

Let $y_t \in [0, 1]$ for $t = 1, \dots, T$ an arbitrary sequence of numbers.

Let the algorithm output be $x_t = x_{t-1}^ = \frac{1}{t-1} \sum_{i=1}^{t-1} y_i$, then*

$$\text{Regret}_T = \sum_{t=1}^T (x_t - y_t)^2 - \sum_{t=1}^T (x_T^* - y_t)^2 \leq 1 + \ln T$$

Lemma

Let x_1^*, x_2^*, \dots be the sequence of predictions produced by FTL. Then for all $u \in R$ (In particular, for $u = x_{T+1}^*$):

$$\begin{aligned} \text{Regret}_T(u) &= \sum_{t=1}^T \left((x_t^* - y_t)^2 - (u - y_t)^2 \right) \\ &\leq \sum_{t=1}^T \left((x_t^* - y_t)^2 - (x_{t+1}^* - y_t)^2 \right) \end{aligned}$$

proof sketch:

Subtract $\sum_{t=1}^T (x_t^* - y_t)^2$ from both sides to get an equivalent claim:

$$\sum_{t=1}^T (x_{t+1}^* - y_t^*)^2 \leq \sum_{t=1}^T (u - y_t)^2$$

The inequality is proven by induction on T .

Sketch of proof of theorem

Using the fact that FTL is $x_t^* = \frac{1}{t-1} \sum_{i=1}^{t-1} y_i$

one can show that $x_{t+1}^* - y_t = \frac{t-1}{t} (x_t^* - y_t)$

and therefor that $(x_t^* - y_t)^2 - (x_{t+1}^* - y_t)^2 = \frac{1}{t} (x_t^* - y_t)^2$

From the fact that $0 \leq x_t^*, y_t \leq 1$ we get that $(x_t^* - y_t)^2 \leq 1$.

From which we obtain $\sum_{t=1}^T ((x_t^* - y_t)^2 - (x_{t+1}^* - y_t)^2) \leq \sum_{t=1}^T \frac{1}{t}$

Combing the last statement with the Lemma concludes the proof of the theorem.